

# Introduction to Computer Graphics


## 9. GPU and Shaders

I-Chen Lin  
National Chiao Tung University

Textbook: E. Angel, D. Shreiner Interactive Computer Graphics, 6th Ed., Pearson  
Ref: D.D. Hearn, M. P. Baker, W. Carithers, Computer Graphics with OpenGL, 4th Ed., Pearson

# The Development of Graphics Cards (consumer-level): Early 90's

- ▶ VGA cards in the early 90's
  - ▶ Just output designated “bitmap”.
  - ▶ Some with 2D acceleration, ex. “Bitblt”
  - ▶ Ex. S3
- ▶ Interactive 3D(or 2.5D) games relied on software rendering.
  - ▶ There were hardware graphics pipelines on workstations, e.g. SGI.



Figures from [https://en.wikipedia.org/wiki/Wolfenstein\\_3D](https://en.wikipedia.org/wiki/Wolfenstein_3D)

# The Development of Graphics Cards (consumer-level): Late 90's

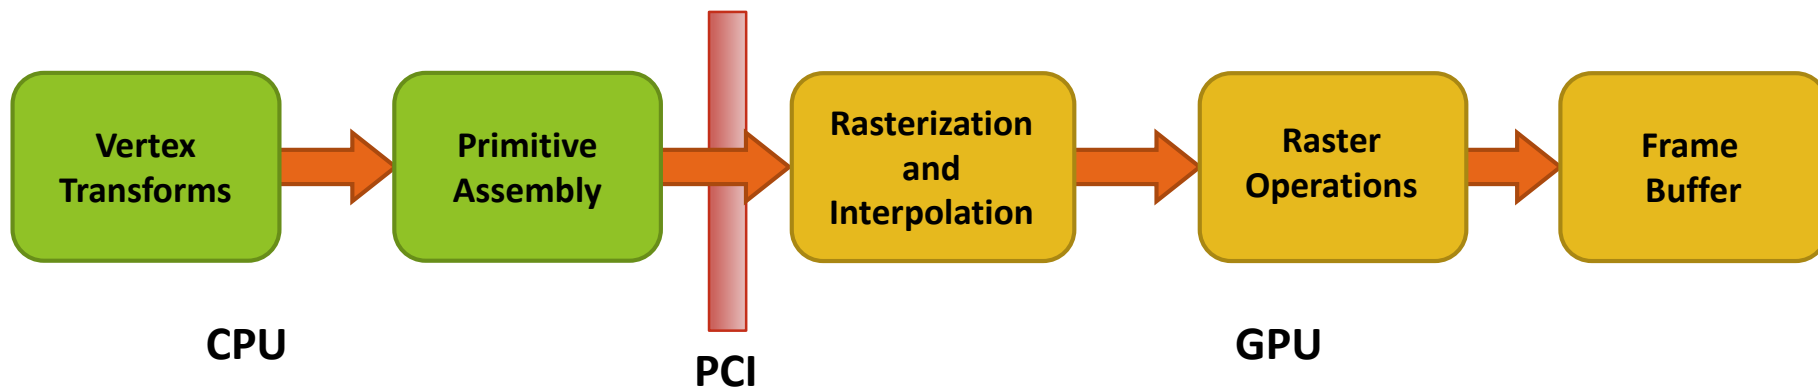
- ▶ 3D accelerators (90's)
  - ▶ Fixed-function pipelines.
  - ▶ E.g. S3, Voodoo, Nvidia, ATI, 3D Labs....
  - ▶ Some of them had to work with a standard VGA card.

# 3Dfx Voodoo (1996)

- ▶ One of the first true 3D game cards
- ▶ Worked by supplementing a standard 2D video card.
- ▶ Did not do vertex transformations (they were evaluated in the CPU)
- ▶ Did texture mapping, z-buffering.



[en.wikipedia.org/wiki/3dfx\\_Interactive](https://en.wikipedia.org/wiki/3dfx_Interactive)



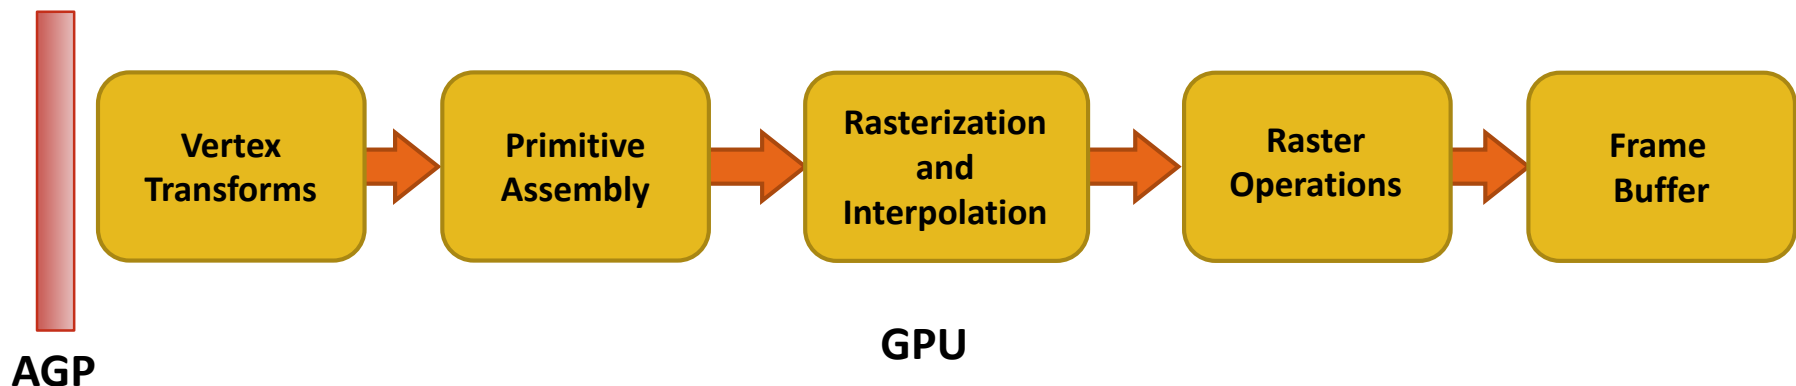
Modified from S. Venkatasubramanian and J. Kider, “Evolution of the Programmable Graphics Pipeline”

# GeForce/Radeon 7500 (1998)

- ▶ Main innovation: shifting the transformation and lighting calculations to the GPU
- ▶ Allowed multi-texturing: giving bump maps, light maps, and others.
- ▶ Faster AGP bus instead of PCI



[en.wikipedia.org/wiki/GeForce\\_256](https://en.wikipedia.org/wiki/GeForce_256)



# The Development of Graphics Cards (consumer-level): after 2001

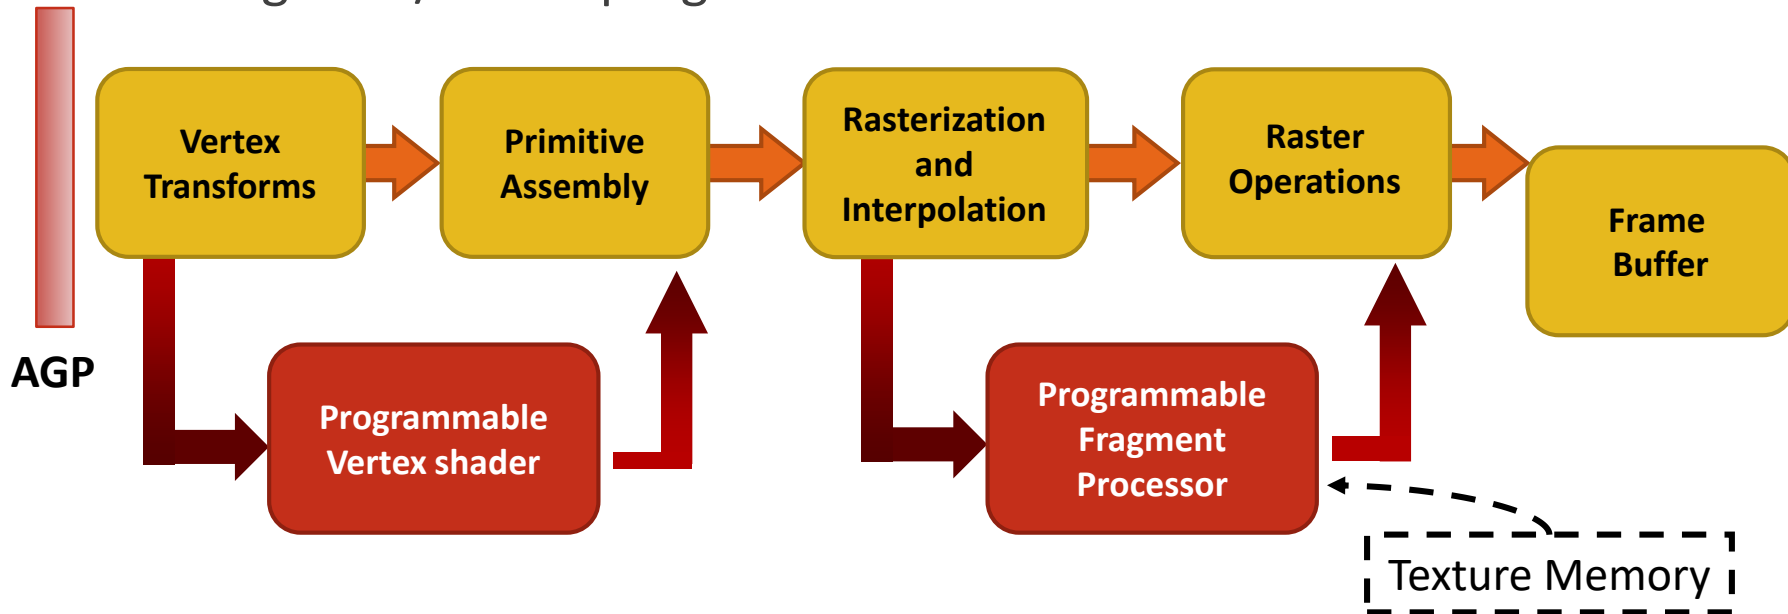
- ▶ Programmable pipelines on GPU
- ▶ GeForce3/Radeon 8500(2001)
  - ▶ Programmable vertex computations: up to 128 instructions
  - ▶ Limited programmable fragment computations: 8-16 instructions



[https://en.wikipedia.org/wiki/GeForce\\_3\\_series](https://en.wikipedia.org/wiki/GeForce_3_series)

# The Development of Graphics Cards (consumer-level): after 2001 (cont.)

- ▶ Radeon 9700/GeForce FX (2002)
  - ▶ the first generation of fully-programmable graphics cards
  - ▶ Different versions have different resource limits on fragment/vertex programs



Modified from S. Venkatasubramanian and J. Kider, "Evolution of the Programmable Graphics Pipeline"

# Evaluation of Graphics Pipeline

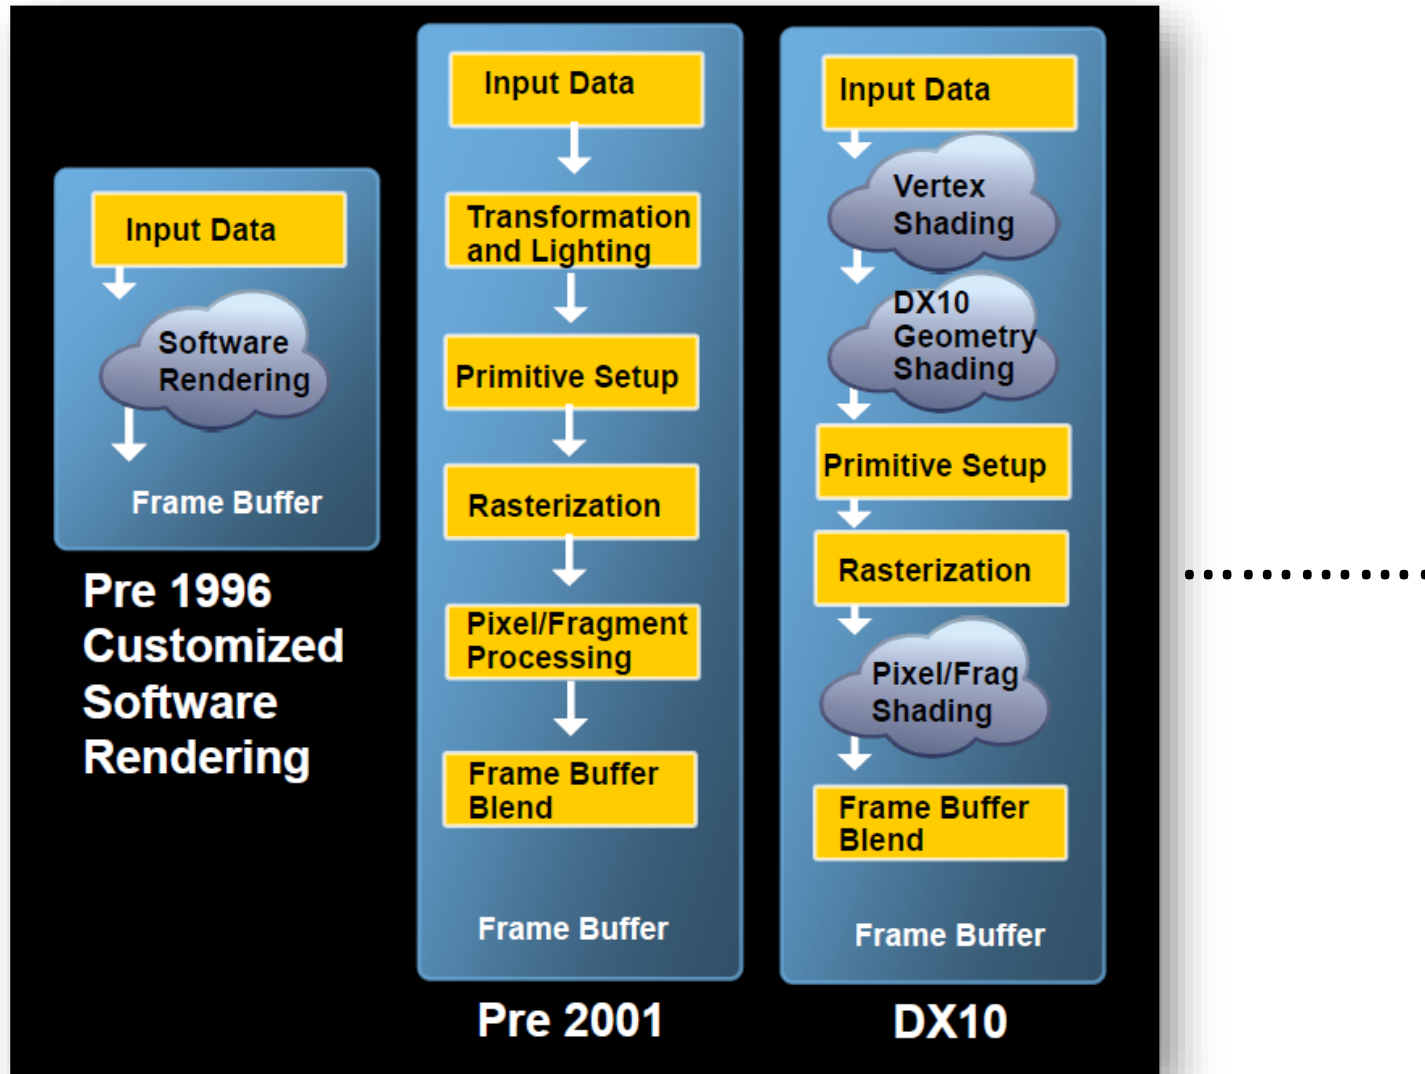


Figure from: M. Houston, “Beyond Programmable Shading Retrospective” slides



# GPU & Shaders : the new age of real-time graphics

- ▶ Programmable pipelines.
- ▶ Supported by high-end commodity cards
  - ▶ NVIDIA, AMD/ATI, etc.




[en.wikipedia.org/wiki/GeForce\\_10\\_series](https://en.wikipedia.org/wiki/GeForce_10_series)



[www.amd.com/zh-hant/products/graphics/radeon-rx-570](http://www.amd.com/zh-hant/products/graphics/radeon-rx-570)

# Why is It So Remarkable?

- ▶ We can do lots of cool stuff in real-time, without overworking the CPU.
  - ▶ Phong Shading
  - ▶ Bump Mapping
  - ▶ Particle Systems
  - ▶ Animation
  - ▶ .....
- ▶ Beyond real-time graphics: GP-GPU, e.g. CUDA, OpenCL (Open Computing Language)
  - ▶ Scientific Data Processing
  - ▶ Computer vision
  - ▶ Deep learning
  - ▶ .....



# Programmable Components

- ▶ Shader: programmable processors.
  - ▶ Replacing fixed-function vertex and fragment processing, and so forth.
- ▶ Shaders:
  - ▶ Vertex shaders
    - ▶ Dealing with per-vertex functions.
    - ▶ We can control the lighting and position of each vertex.
  - ▶ Fragment shaders
    - ▶ Dealing with per-pixel functions.
    - ▶ We can control the color of each pixel by user-defined programs.
  - ▶ Geometry shaders (DirectX 10, SM 4+)
  - ▶ New shaders (hull, domain) in DirectX11, SM5

# Programmable Components (cont.)

- ▶ Software Support

- ▶ Direct X 8 , 9, 10, 11, 12, ...

- ▶ OpenGL Extensions

- ▶ OpenGL Shading Language (GLSL)

- ▶ OpenGL for Embedded Systems (OpenGL ES)

- ▶ *Cg (C for Graphics)*

- ▶ Metal Shading Language (by Apple)

- ▶ .....

# Vertex Shaders

- ▶ Per-vertex calculations performed here
  - ▶ Without knowledge about other vertices (parallelism)
- ▶ Your program take responsibility for:
  - ▶ Vertex transformation
  - ▶ Normal transformation
  - ▶ (Per-Vertex) Lighting
  - ▶ Color material application and color clamping
  - ▶ Texture coordinate generation


# Vertex Shader Applications

- ▶ We can control movement with uniform variables and vertex attributes
  - ▶ Time
  - ▶ Velocity
  - ▶ Gravity
- ▶ Moving vertices
  - ▶ Morphing
  - ▶ Wave motion
  - ▶ .....
- ▶ Lighting
  - ▶ More realistic models
  - ▶ Cartoon shaders

# Applications: Wave Motion Vertex Shader



```
uniform float time;  
uniform float xs, zs;  
void main()  
{  
    float s;  
    s = 1.0 + 0.1*sin(xs*time)*sin(zs*time);  
    gl_Vertex.y = s*gl_Vertex.y;  
    gl_Position =  
    gl_ModelViewProjectionMatrix*gl_Vertex;  
}
```

Uniform: passing parameters to vertex and fragment shaders.



# Applications: Particle Systems

```
uniform vec3 init_vel;  
uniform float g, m, t;  
void main()  
{  
    vec3 object_pos;  
    object_pos.x = gl_Vertex.x + vel.x*t;  
    object_pos.y = gl_Vertex.y + vel.y*t  
    + g/(2.0*m)*t*t;  
    object_pos.z = gl_Vertex.z + vel.z*t;  
    gl_Position =  
    gl_ModelViewProjectionMatrix*  
    vec4(object_pos,1);  
}
```



Uniform: passing parameters to vertex and fragment shaders.

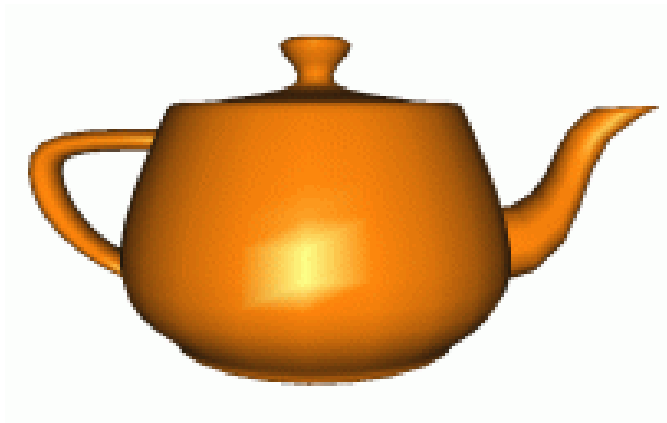


# Fragment Shaders

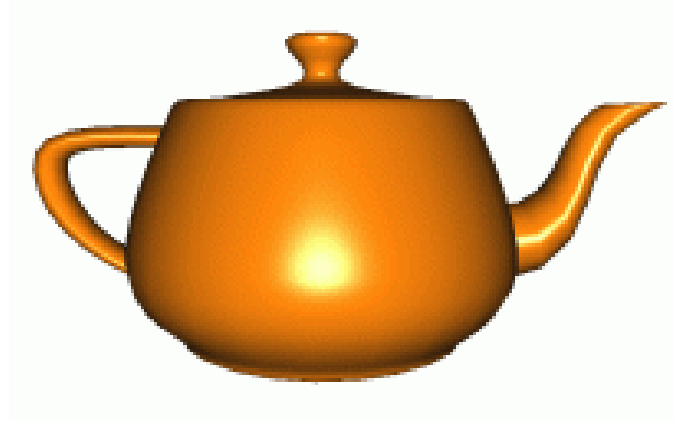
- ▶ What is a fragment?
  - ▶ Cg Tutorial says: “You can think of a fragment as a ‘potential pixel’”
- ▶ Perform per-pixel calculations
  - ▶ Without knowledge about other fragments (parallelism)
- ▶ Your program’s responsibilities:
  - ▶ Operations on interpolated values
  - ▶ Texture access and application
  - ▶ Other functions: fog, color lookup, etc.

# Fragment Shader Applications

(Per-pixel) Phong shading



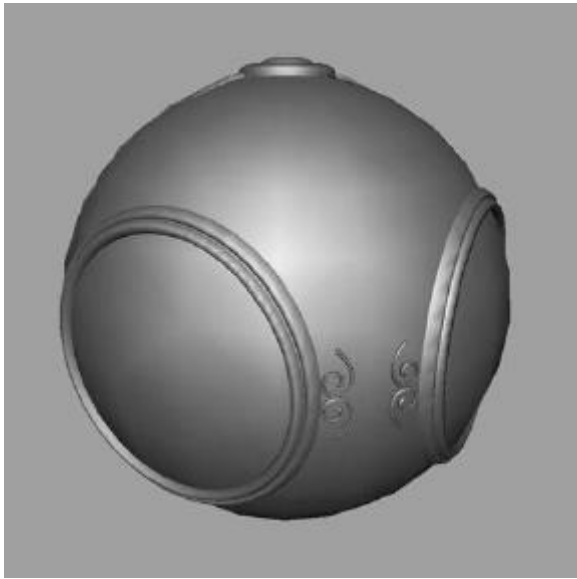
Per-vertex lighting



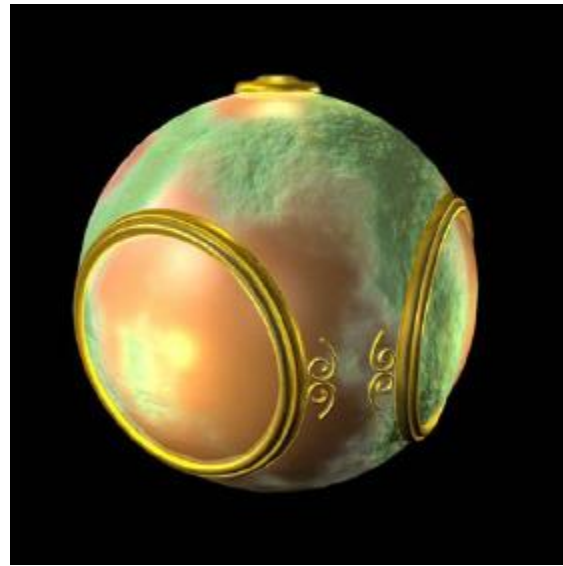
Per-fragment lighting

Figures from <http://www.lighthouse3d.com/opengl/glsl/>

# Fragment Shader Applications





smooth shading



bump mapping

# Bump Mapping

- ▶ Perturb normal for each fragment
- ▶ Store perturbation as textures



# Toon Shading

Note: varying, communicating between vertex and fragment

- The vertex shader then becomes:

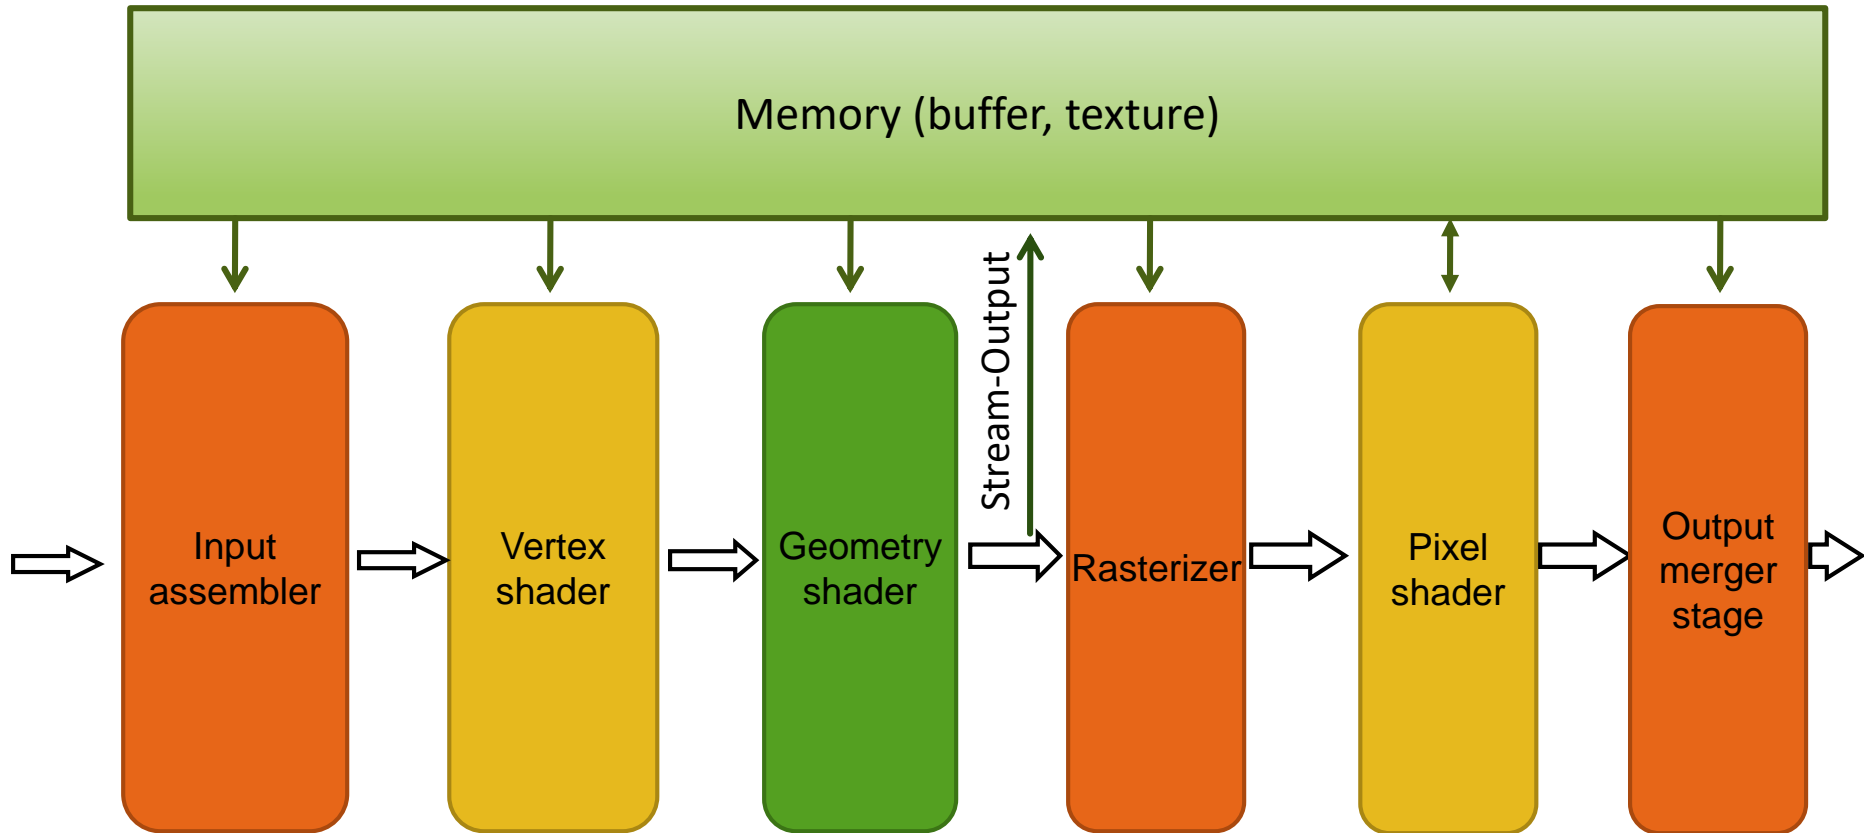
```
varying vec3 normal;  
void main() {  
    normal = gl_NormalMatrix * gl_Normal;  
    gl_Position = ftransform(); }  
}
```

- The pixel shader becomes

```
varying vec3 normal;  
void main() {  
    float intensity; vec4 color;  
    vec3 n = normalize(normal);  
    intensity = dot(vec3(gl_LightSource[0].position),n);  
    if (intensity > 0.95) color = vec4(1.0,0.5,0.5,1.0);  
    else if (intensity > 0.5) color = vec4(0.6,0.3,0.3,1.0);  
    else if (intensity > 0.25) color = vec4(0.4,0.2,0.2,1.0);  
    else color = vec4(0.2,0.1,0.1,1.0);  
    gl_FragColor = color; }  
}
```



# With the Geometry Shader



Direct3D 10 pipeline stage from MSDN of Microsoft

# D3D 10 Pipeline

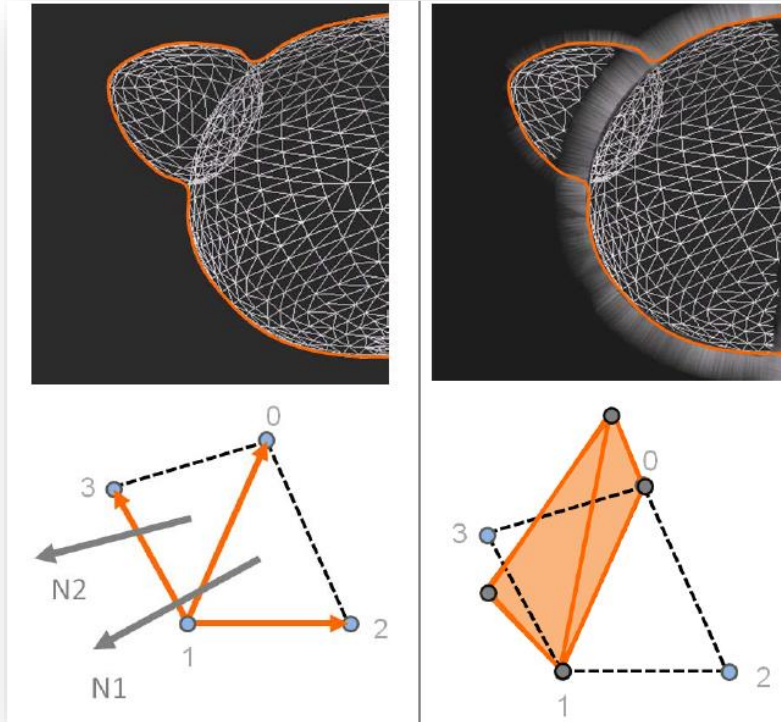
- ▶ **Input assembler:** supplies data (triangles, lines and points) to the pipeline.
- ▶ **Vertex shader:** processes vertices, such as transformations, skinning, and lighting.
- ▶ **Geometry shader:** processes entire primitives.
  - ▶ 3 vertices: a triangle, 2 vertices: a line, or 1 vertex: a point.
  - ▶ The Geometry shader supports limited geometry amplification and de-amplification. (discard the primitive, or emit one or more new primitives)
  - ▶ E.g. Subdivision, point -> billboard, silhouette edge -> fur, etc.
- ▶ **Stream-output stage:**
  - ▶ Data can be streamed out and/or passed into the rasterizer. Data streamed out to memory can be recirculated back into the pipeline as input data or read-back from the CPU.

## D3D 10 Pipeline (cont.)


- ▶ **Rasterizer:** clips primitives, prepares primitives for the pixel shader and determines how to invoke pixel shaders.
- ▶ **Pixel shader:** receives interpolated data for a primitive and generates per-pixel data, such as color.
- ▶ **Output-merger stage:**
  - ▶ combines various types of output data (pixel shader values, depth and stencil information) with the contents of the render target and depth/stencil buffers to generate the final pipeline result.



# D3D 10 Pipeline (cont.)



Figures from NVIDIA DirectX10 SDK Doc:  
Fur (using Shells and Fins)



# D3D 11 Pipeline

- ▶ In D3D10, the Geometry shader may subdivide the surfaces by multiple passes.
- ▶ D3D11 improves the tessellation ability by three new stages: hull shader, tessellator, domain shader.
- ▶ The tessellated patches can still be applied to geometry shaders. E.g. point -> billboard, silhouette edge -> fur, etc.

# Tessellation Pipeline

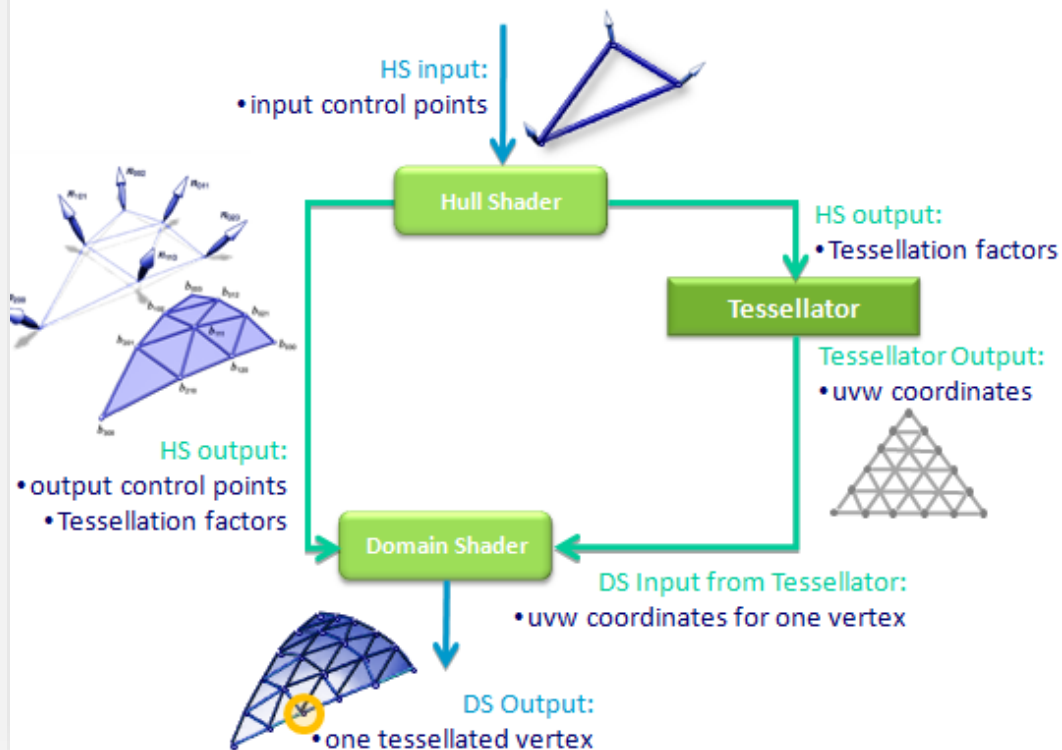


Figure from:  
[developer.download.nvidia.com/presentations/2009/GDC/GDC09\\_D3D11Tessellation.pdf](http://developer.download.nvidia.com/presentations/2009/GDC/GDC09_D3D11Tessellation.pdf)

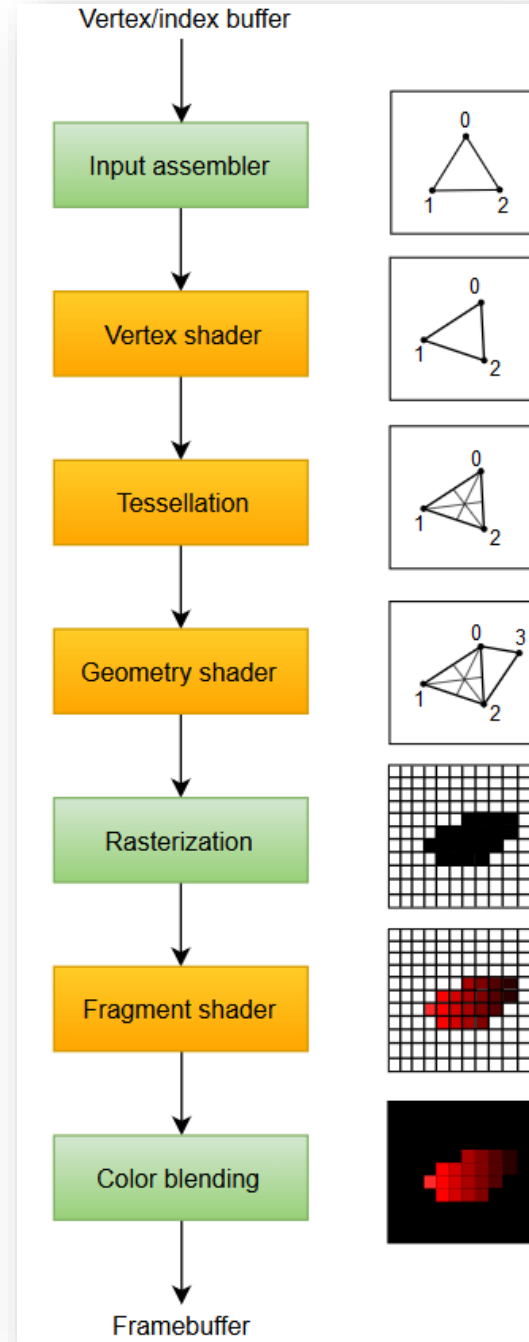


Figure from: [vulkan-tutorial.com/Drawing\\_a\\_triangle/Graphics\\_pipeline\\_basics/Introduction](http://vulkan-tutorial.com/Drawing_a_triangle/Graphics_pipeline_basics/Introduction)

# D3D 11 Tessellation



Model refinement



Base  
Model




Bump  
Mapping



Displacement  
Mapping

Image courtesy of [www.chromosphere.com](http://www.chromosphere.com)



Tessellation with displacement mapping

End of Chapter 9